AuditMind — Multi-Scenario Demo & Presentation Playbook

Food-Safety Compliance Agent Powered by Hindsight Long-Term Memory & Band AI

Executive Briefing: AuditMind solves restaurant chain audit failures by converting fragmented paper checklists and chat messages into a persistent, learning memory bank. It operates in **Band Desktop** and local CLI, backed by **Vectorize Hindsight** and **Groq LLM**. This playbook equips presenters with 5 structured scenarios, before/after contrast scripts, and live defense answers.

1. Presentation Blueprint & Stagecraft (3-Minute Flow)

Judges evaluate 5 themes: **Innovation (30%)**, **Hindsight Memory (25%)**, **Technical Implementation (20%)**, **User Experience (15%)**, and **Real-world Impact (10%)**. Follow this strict timing budget to maximize marks.

Time	Phase	Screen Action	Key Talking Point
0:00 - 0:30	The Real Problem	Show fragmented WhatsApp/paper logs vs. restaurant penalties.	'Restaurants fail audits for repeat causes: the same fridge, the same supplier. AI today forgets.'
0:30 - 1:00	The 'Before' Shot	Run CLI with `/memory off`. Ask audit prep.	'Without memory, the model spits out a textbook checklist with zero outlet context.'
1:00 - 1:50	The 'After' Shot (Memory Star)	Toggle `/memory on`. Ask same question. Point out colored traces.	'Watch [hindsight:REFLECT]. It recalls past breaches, failed fixes, and flags the pest gap.'
1:50 - 2:30	Live Ingestion & Learning	Log warm delivery as Arjun. Ask weekly briefing as Meera.	'It connects deliveries across 3 sites and spots that night shift is accepting warm poultry.'
2:30 - 3:00	Impact & Wrap-Up	Show architecture slide or Band Desktop room view.	'Memory transforms a disposable chatbot into a proactive corporate compliance auditor.'

Scenario 1: Audit Readiness & Historical Risk Ranking (Core Pitch)

Objective: Demonstrate the immediate, undeniable difference between a stateless LLM and Hindsight memory.

Step	Prompt & Speaker	Agent Response & System Behavior

1. Before	**Speaker:** Meera
	`/memory off`
	`Is Koramangala ready for Friday's audit?`
		Generic Checklist: Gives textbook advice: *'Check fridge temp ≤5°C, verify employee hairnets, ensure daily cleaning logs are signed.'* Completely unaware of Koramangala's real history.

2. After	**Speaker:** Meera
	`/memory on`
	`Is Koramangala ready for Friday's audit?`
		Memory-Grounded Intelligence:
		• Triggers `[hindsight:REFLECT]`.
		• Identifies 8 past morning breaches in walk-in fridge.
		• Recalls that replacing thermometer **failed**, but moving marination to reach-in 2 **worked**.
		• Flags the 6-day pest log gap from Ayudha Puja week (a repeat finding from March client audit).
		• Prescribes exact 7-day temperature verification logs to hand to the auditor.

Scenario 2: The 'Fix That Failed' vs 'Fix That Worked'

Objective: Show that fix outcomes are first-class memories. The agent stops recommending fixes that already failed.

Interaction	Details & Code-Backed Flow
Speaker & Context	Arjun Menon (Koramangala Manager) reporting ongoing walk-in fridge issues.
Manager Message	`Walk-in fridge is reading 7.2°C this morning. What should we do?`
Agent Reasoning	1. Recalls `fix-outcome` record: On Oct 4, replacing thermometer failed (readings remained 6.5–7.5°C). 2. Recalls successful fix: On Oct 17, moving morning chicken marination out of walk-in dropped temps to 3–4°C. 3. Rejects recommending a thermometer replacement.
Agent Output	*'Logged breach (7.2°C vs ≤5°C limit). Note: replacing the thermometer previously did NOT work here. Check whether morning prep marination trays were loaded into the walk-in with the door left open. Moving marination to reach-in fridge 2 solved this in October.'*

Scenario 3: Cross-Site Pattern & Supplier Intelligence

Objective: Highlight the architectural decision of having **one shared memory bank** across all outlets rather than isolated silos.

Trigger Event	Cross-Outlet Correlation & Root Cause Discovery
New Log at Koramangala	Arjun logs: `FreshCo chicken delivery just came in at 7.8°C, Imran took it on night shift.`
Tool Invocation	`log_delivery(outlet='Koramangala', supplier='FreshCo Poultry', temp_c=7.8, accepted=true, shift='night')`
Cross-Site Insight	Hindsight recalls deliveries across Indiranagar, Koramangala, and HSR Layout: • 5 warm deliveries from FreshCo (6.9°C to 8.4°C). • 4 accepted by night-shift staff (driver in a hurry); 1 rejected by day shift. • Correlates root cause: FreshCo changed delivery route to post-10 PM; night staff lack receiving discipline.
Executive Briefing Response	Meera asks: `weekly risk briefing` Agent synthesizes: *'Primary group risk is FreshCo Poultry late-night deliveries across all 3 sites. Immediate action: review supplier route SLA and retrain night receivers.'*

Scenario 4: Temporal Policy Versioning & Evidence Gaps

Objective: Demonstrate Gate 4 compliance — the agent respects time, policy versions, and refuses to guess missing evidence.

Compliance Rule	Prompt Tested	Observed Compliant Output
Non-Retroactivity (Policy v1 vs v2)	`Did the Feb 20 Whitefield delivery violate the April receiving documentation rule?`	**No retroactive penalty:** Agent states REQ-003 was active in Feb. REQ-004 took effect April 1, 2026. The delivery had a temp breach, but cannot be penalized under unannounced documentation rules.
Action Completed vs. Verified	`Is the Indiranagar temperature finding closed after the June 21 fix?`	**Distinguishes completion:** Agent notes the supervisor sign-off was assigned, but remained overdue until Oct. Only on Oct 18 did 7-day electronic verification close the finding.
Preserving Uncertainty	`Does the Whitefield Sept record prove chicken batch compliance?`	**Refuses to hallucinate:** Temperature was 3.9°C (pass), but batch number was blank. Supervisor claimed invoice had it, but invoice wasn't attached. Agent flags as incomplete gap.

Scenario 5: Natural Conversational UX & Batched Intake

Objective: Show high UX score (15%) by demonstrating human-like communication and zero repetitive multi-turn badgering.

Input Case	User Input	Agent Behavior & Output
Casual Greeting	`hi` or `good morning`	Warm, natural 2-line greeting introducing AuditMind; triggers no unnecessary tool calls; invites user needs directly.
Vague Incident	`fridge is warm at Koramangala`	**Batches all questions into one:** *'Quick check before I log this:* *1. Which fridge is it (walk-in, reach-in)?* *2. What was the exact temperature reading (°C)?* *3. Who recorded it and which shift?'*
Acknowledgement	`thanks, got it`	Polite, brief sign-off (*'Happy to help! Let me know if anything else comes up.'*) with zero robotic filler.

Judge Defense & Technical Q&A; Cheat Sheet

Likely Judge Question	Winning Answer Grounded in Architecture
Why use Hindsight instead of standard RAG with pgvector?	Traditional RAG does simple semantic cosine matching. It has no temporal graph, cannot detect that a fix *failed* vs *succeeded*, and cannot connect that event B contradicted event A. Hindsight provides `reflect` and dynamic disposition (skepticism: 4) to catch contradictions.
Why not give each outlet its own memory bank?	Siloed banks prevent group-level intelligence. The FreshCo poultry defect pattern only became visible because deliveries from all 3 outlets fed into a single restaurant group bank with site metadata.
How does the system handle LLM function calling errors?	In `src/llm.ts`, we built custom retry logic that catches HTTP 400 `tool_use_failed` and 429 rate limits, injects corrective system instructions, and gracefully falls back to direct responses if retries expire.
Is this using real patient or restaurant data?	No. All 62 longitudinal records and 42 synthetic daily logs are fictional and generated with deterministic seeds. All compliance rules are marked internal illustrative standards, adhering to privacy rules.